

# Comparative Evaluation of Segmentation and Object Detection Models for Pallet Recognition

Aksheya Kannan Subramanian

Michael Acquah

Faten Alshohatee

CIS 583 – Deep Learning

Each slide in this presentation has an accompanying audio recording. Feel free to play them manually or view the presentation in SlideShow mode for the full experience.






# Motivation



AUTONOMOUS  
WAREHOUSES  
REQUIRE RELIABLE  
PERCEPTION FOR SAFE  
ROBOT NAVIGATION



TWO KEY TASKS:  
**GROUND  
SEGMENTATION AND  
PALLET DETECTION**



TRADITIONAL CV  
STRUGGLES WITH  
LIGHTING, SHADOWS,  
REFLECTIONS



DEEP LEARNING  
PROVIDES STRONGER  
ROBUSTNESS FOR  
REAL INDUSTRIAL  
ENVIRONMENTS



GOAL: COMPARE  
SEGMENTATION AND  
DETECTION MODELS  
FOR WAREHOUSE  
DEPLOYMENT





# Research Questions

How do **U-Net** and **DeepLabV3+** compare for warehouse ground segmentation?

How do **YOLOv8** and **EfficientDet-D0** differ in pallet detection performance?

What are the key trade-offs in **accuracy**, **latency/FPS**, **confidence calibration**, and **model size** for real-time deployment?



# Dataset: EuroPalletSeg



Real indoor warehouse scenes with pallets, shelving, and concrete flooring



Challenging visual conditions: **shadows, reflections, partial occlusions**



**Ground segmentation masks**



**YOLO-format pallet bounding boxes**




Separate train/validation splits with **585 validation images** for each task





# Dataset Creation



# Preprocessing

## Segmentation Pipeline

- Resize images to **256×256**
- Apply model-specific normalization
  - U-Net → scale to **[0,1]**
  - DeepLabV3+ → **ImageNet mean/std**
- Resize masks with **nearest-neighbor** interpolation
- Apply binary **thresholding (>127)** to clean the mask

## Detection Pipeline

- Resize images to **512×512**
- Convert YOLO labels from **normalized → pixel coordinates**
- Use absolute coordinates for IoU + evaluation consistency

## EfficientDet Postprocessing

- EfficientDet-D0 outputs many raw predictions
- Only the **top 5 highest-confidence** boxes per image were kept for comparison + visualization



# Models Evaluated

## Segmentation Models

### U-Net — 209 MB

Encoder-decoder with skip connections; high spatial precision

### DeepLabV3+ (MobileNetV3) — 18 MB

Atrous Spatial Pyramid Pooling; lightweight + efficient

## Detection Models

### YOLOv8 — 42 MB

One-stage detector with fast, confident predictions

### EfficientDet-D0 — 15 MB

Compact detector using EfficientNet + BiFPN





# U-Net Overview

## Architecture

- Classical **encoder–decoder** structure
- Uses **skip connections** to recover spatial detail lost during downsampling

## Strengths

- Produces **sharp, accurate segmentation masks**
- Well-suited for tasks requiring precise boundary reconstruction
- **Large model size (209 MB)** → less efficient for embedded or real-time deployment





# DeepLabV3+ Overview

## Architecture

- **Atrous Spatial Pyramid Pooling (ASPP)** for multi-scale context
- **MobileNetV3 encoder** for efficient feature extraction
- Decoder refines segmentation boundaries

## Strengths

- **Lightweight (18 MB)** and fast
- Captures both local detail and global context
- Well-suited for **resource-constrained** or real-time robotics





# YOLOv8 Overview

## Architecture

- **One-stage detector** with a modern C2f backbone
- Decoupled detection head improves box + class predictions

## Strengths

- **Fast inference** and low latency on CPU
- **High confidence scores** and stable predictions
- **Robust pallet detection**, even under reflections and occlusions





# EfficientDet-D0 Overview

## Architecture

- **EfficientNet-B0 backbone** for compact, high-quality feature extraction
- **BiFPN** (Bidirectional Feature Pyramid Network) for efficient multi-scale fusion
- Designed with compound scaling for balanced depth, width, and resolution

## Strengths

- **Very small model size (~15 MB)**
- Highly **efficient and lightweight**, ideal for low-power or embedded devices
- Optimized for speed while maintaining reasonable accuracy on general object detection tasks





# Evaluation Metrics

## Segmentation Metrics

**IoU (Intersection over Union)** – overlap quality

**Dice Coefficient** – foreground accuracy

**Pixel Accuracy** – overall correctness

**Precision / Recall** – false positives vs. false negatives

## Detection Metrics

**mAP@0.5** – bounding box quality


**FPS** – inference speed

**Latency** – per-image processing time

**Confidence Scores** – model certainty in detections



# Segmentation Results




Model	IoU	Dice
U-Net	0.951	0.973
DeepLabV3+	0.921	0.973



# Detection Results

Model	mAP@0.5	Latency (ms)	FPS
EfficientDet-D0	0.000	158.55	6.30
YOLOv8	0.0014	95.44	10.47





# Qualitative Results



**U-Net:** produces sharper, more detailed ground masks



**DeepLabV3+:** smooth and coherent masks, slightly less precise at boundaries



**YOLOv8:** consistent and reliable pallet bounding boxes



**EfficientDet-D0:** weak detections with very low confidence, often missing pallets



# Limitations

## 📌 Experimental Constraints

CPU-only evaluation, limiting achievable speed

Restricted hyperparameter tuning due to compute constraints

## 📌 Model-Specific Challenges

EfficientDet-D0 required confidence calibration, which was not fully explored

## 📌 Scope Limitations

No multi-task or joint segmentation–detection model evaluated



# Future Work

---

Hyperparameter tuning and calibration for  
**EfficientDet-D0**

---

Explore next-generation detectors such as  
**YOLOv9** or transformer-based models

---

Develop a **multi-task model** combining  
segmentation + detection in one network

---

Test models on **edge hardware** (e.g.,  
Jetson, Raspberry Pi, industrial robotics  
hardware)

---

Evaluate performance after optimization  
(quantization, pruning)





# Conclusion



**U-Net** produced the most accurate segmentation, while **DeepLabV3+** was nearly as effective but much more efficient.



**YOLOv8** was the most practical and reliable for pallet detection, while **EfficientDet-D0** needed more tuning despite being lightweight.



Our findings suggest using **U-Net** or **DeepLabV3+** based on compute resources and deploying **YOLOv8** for reliable pallet detection.





**THANK YOU!**

